


Build Your Own Router

Computer Networks
Course Project

Project Overview

- Implement a simple router in a single router topology with **static routing table** (forwarding table).
- Your router will receive **raw Ethernet frames**, and handle/forward packets in correct logic.



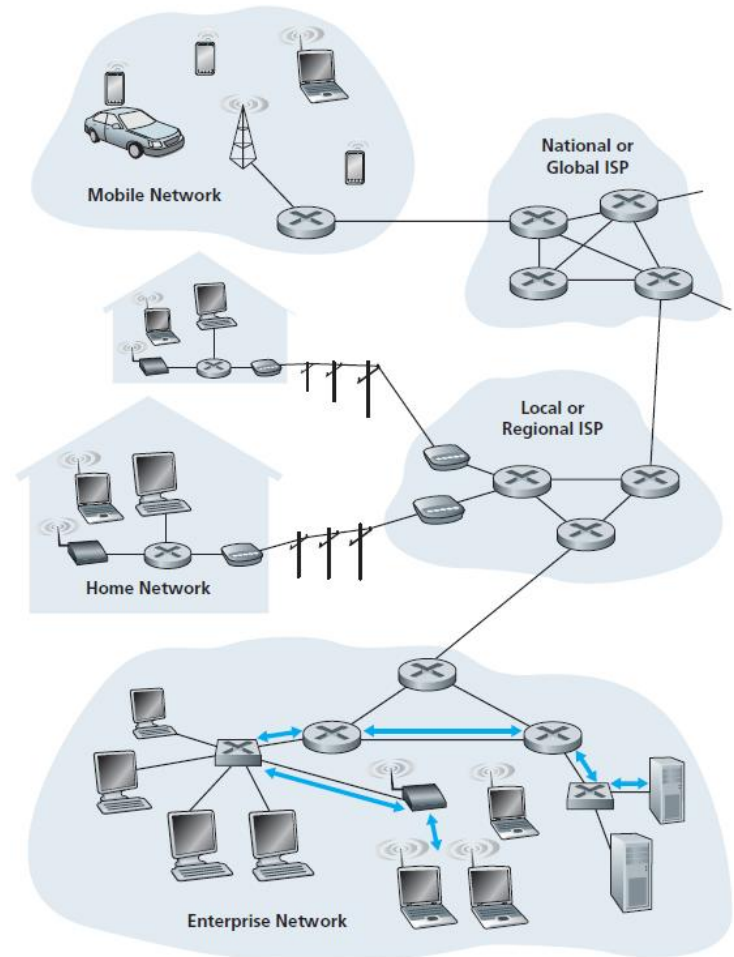
LESSONS TO LEARN

What about Our Router?

- **Forwarding:** move packets from router's input to appropriate router output
 - Load predefined routing table
 - Look up matching entry in routing table
- Handle ICMP
- Handle ARP request / reply
- ~~**Routing:** determine route taken by packets from source to dest.~~

Link Layer


- Transfer internet layer datagram from **node to adjacent node** over a link
- Encapsulate internet layer **datagram into frame**, add header & trailer
- Use "**MAC**" address in frame headers for source, dest



MAC address v.s. IP address

- 32-bit IP address:
 - network-layer address
 - used to get datagram from src to dest IP subnet
- 48-bit MAC (Ethernet) address:
 - link-layer address for network interfaces
 - get frame from **one interface to another physically-connected interface (same subnet)**
 - Broadcast address: "FF:FF:FF:FF:FF:FF"

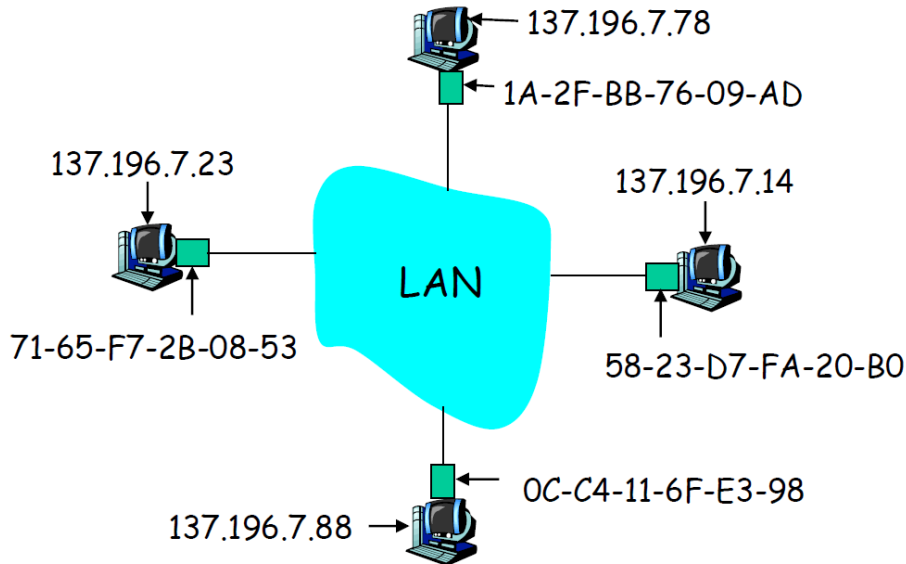
Ethernet Frame



- Encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**
 - preamble and crc are handled transparently in this project


ARP: Address Resolution Protocol

Question: How to get
MAC address of B from
B's IP address




- ARP request: request IP-MAC mapping of next hop interface
(send to broadcast address)
- ARP reply: send IP-MAC of current interface
- ARP cache: **IP - MAC mapping** for nodes
(timeout after a 30s)

ARP: Address Resolution Protocol



Walkthrough: send datagram from A to B through R.

Internet Protocol (IPv4)




- Delivering packets from src to dest based on IP address
- Header **checksum**
- Decrement **TTL**

Routing Table

- Destination & Netmask: subnet **network ID**
- Gateway: next hop IP address to destination
- Interface: name of network interface card connected to gateway
- Metric: routing metric of path to destination
(omitted in this project)

Longest Prefix Match


| Destination | Netmask | Gateway | Interface |
|-------------|---------------|------------|-----------|
| 0.0.0.0 | 0.0.0.0 | 10.0.1.100 | eth3 |
| 192.168.2.2 | 255.255.255.0 | * | eth1 |
| 172.64.3.10 | 255.255.0.0 | * | eth2 |



Where should packet dest at "192.168.1.1" go?


Longest Prefix Match

| Destination | Netmask | Gateway | Interface |
|-------------|---------------|------------|-----------|
| 0.0.0.0 | 0.0.0.0 | 10.0.1.100 | eth3 |
| 192.168.2.2 | 255.255.255.0 | * | eth1 |
| 172.64.3.10 | 255.255.0.0 | * | eth2 |



Where should packet dest at "192.168.2.1" go?

ICMP: Internet Control Message Protocol



- Used by hosts & routers to communicate network-level information
 - error reporting (unreachable host, network, port, protocol)
 - echo request/reply
- Sent as **IPv4 payload**
- Content: Internet header + 8 bytes of original datagram

Ping & Traceroute

- **Ping:** ICMP echo request/reply
 - reply with TTL = 64

Ping & Traceroute

- **Ping:** ICMP echo request/reply
 - reply with TTL = 64
- **Traceroute:** displaying possible routes and RTT in IP network.
 - sends UDP segments with TTL = 1,2,3, ... with unlikely port number
 - when nodes receive datagram with TTL=0, return ICMP "Time Exceeded"
 - destination returns ICMP "Port Unreachable"


Useful Materials

- IPv4:
 - RFC 791: <https://tools.ietf.org/html/rfc791>
 - Text Book: section 4.4.1, 4.4.2
- ICMP:
 - RFC 792: <https://tools.ietf.org/html/rfc792>
 - Text Book: section 4.4.3
- ARP:
 - RFC 826: <https://tools.ietf.org/html/rfc826>
 - Text Book: section 5.4.1

RESULTS TO SHOW

Expected Behaviors

- **ping** from client to any server & router interfaces
- **traceroute** from client to any server & router interfaces
- **wget** files from server
- Update ARP cache table



Ping

- ping from client to any http servers

```
mininet> client ping server1
PING 192.168.2.2 (192.168.2.2) 56(84) bytes of data.
64 bytes from 192.168.2.2: icmp_seq=1 ttl=63 time=1293 ms
64 bytes from 192.168.2.2: icmp_seq=2 ttl=63 time=312 ms
64 bytes from 192.168.2.2: icmp_seq=3 ttl=63 time=50.3 ms
64 bytes from 192.168.2.2: icmp_seq=4 ttl=63 time=29.3 ms
^C
--- 192.168.2.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3025ms
rtt min/avg/max/mdev = 29.388/421.625/1293.814/515.809 ms, pipe 2
```

Ping

- ping wrong IP address

```
mininet> client ping 192.168.2.3
PING 192.168.2.3 (192.168.2.3) 56(84) bytes of data.
From 10.0.1.1 icmp_seq=1 Destination Host Unreachable
From 10.0.1.1 icmp_seq=2 Destination Host Unreachable
From 10.0.1.1 icmp_seq=3 Destination Host Unreachable
From 10.0.1.1 icmp_seq=4 Destination Host Unreachable
From 10.0.1.1 icmp_seq=5 Destination Host Unreachable
^C
--- 192.168.2.3 ping statistics ---
6 packets transmitted, 0 received, +5 errors, 100% packet loss, time 5100ms
pipe 5
```

Traceroute

- traceroute to any http server

```
mininet> client traceroute server1
traceroute to 192.168.2.2 (192.168.2.2), 30 hops max, 60 byte packets
 1  10.0.1.1 (10.0.1.1)  12.806 ms  13.727 ms  14.505 ms
 2  192.168.2.2 (192.168.2.2)  99.179 ms  104.646 ms  106.050 ms
```

- traceroute to router interfaces

```
mininet> client traceroute 192.168.2.1
traceroute to 192.168.2.1 (192.168.2.1), 30 hops max, 60 byte packets
 1  10.0.1.1 (10.0.1.1)  796.441 ms  839.769 ms  839.814 ms
```

File Downloading

- wget from any of servers

```
mininet> client wget http://192.168.2.2/tmp
--2021-11-09 17:18:15--  http://192.168.2.2/tmp
Connecting to 192.168.2.2:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 10240 (10K) [application/octet-stream]
Saving to: 'tmp'

tmp                                100%[=====>]  10.00K  --.-KB/s    in 0.06s

2021-11-09 17:18:17 (163 KB/s) - 'tmp' saved [10240/10240]
```

ARP Cache table

- “pingall”

```
Every 1.0s: ./show-arp.py Tue Nov 9 17:20:53 2021
```

| MAC | IP | AGE | VALID |
|-------------------|-------------|-----------|-------|
| 9a:45:f8:6e:47:2b | 192.168.2.2 | 2 seconds | 1 |
| 52:b1:d0:a7:a9:c9 | 10.0.1.100 | 1 seconds | 1 |
| 82:10:d6:23:40:ad | 172.64.3.10 | 0 seconds | 1 |


- 30 seconds later

```
Every 1.0s: ./show-arp.py Tue Nov 9 17:21:24 2021
```

| MAC | IP | AGE | VALID |
|-----|----|-----|-------|
|-----|----|-----|-------|

CODE TO IMPLEMENT

Code Structure



Key Methods

- **NEED TO IMPLEMENT**
- Method that receives a raw Ethernet frame (simple-router.hpp/cpp):

```
/**  
 * This method is called each time the router receives a packet on  
 * the interface. The packet buffer \p packet and the receiving  
 * interface \p inIface are passed in as parameters.  
 */  
void  
SimpleRouter::handlePacket(const Buffer& packet, const std::string& inIface);
```

Key Methods

- **IMPLEMENTED**
- Method to send raw Ethernet frames (simple-router.hpp/cpp):

```
/**  
 * Call this method to send packet \p packet from the router on interface \p outIface  
 */  
void  
SimpleRouter::sendPacket(const Buffer& packet, const std::string& outIface);
```

Key Methods

- **NEED TO IMPLEMENT**
- Method to handle ARP cache events (arp-cache.hpp/cpp):

```
/**  
 * This method gets called every second. For each request sent out,  
 * you should keep checking whether to resend a request or remove it.  
 */  
void  
ArpCache::periodicCheckArpRequestsAndCacheEntries();
```

Key Methods

- **NEED TO IMPLEMENT**
- Method to lookup entry in the routing table (routing-table.hpp/cpp):

```
/**  
 * This method should lookup a proper entry in the routing table  
 * using "longest-prefix match" algorithm  
 */  
RoutingTableEntry  
RoutingTable::lookup(uint32_t ip) const;
```

ISSUES TO NOTICE

环境配置

- 作业文档中提供了详细的环境配置指南
- 为了方便同学们快速配置环境，我们也提供了虚拟机文件
 - Windows VMware:
<https://cloud.tsinghua.edu.cn/d/05b4618dad6046b49438/>
 - Mac(Apple Silicon) UTM:
<https://cloud.tsinghua.edu.cn/f/d09f8e35213540a5a2cf/>
- 虚拟机的账号和密码都是 **router**

Some Important Issues

- Grading (up to 105%)
 - Router Implementation (85%=45%public + 40%private)
 - Ping tests (50%)
 - Traceroute tests (20%)
 - File Downloading tests (15%)
 - Project Report + Code Quality (20%)
- Individual work

Some Important Issues

- Submission
 - Source code (*"make tarball"*)
 - Do not modify **existing data structures**
 - Report: no longer than **THREE** pages
- Evaluation
 - ping, traceroute, file downloading, details in **project spec**
 - Code quality
 - Project report

补交规则

- 正常提交: 在 DDL 前提交作业正常计分
- 最迟提交期限和惩罚
 - 每次作业最迟在 DDL 后一周 (7天) 内提交, 超出此期限**一律拒收**
 - 未超出此拒收期限的迟交作业得分*0.8.
- 宽限期: 全部大作业共享 7 天宽限期。迟交**累计不超过**该期限的免于扣分惩罚

Some Important Issues

- Start up: today, after class
- Deadline: 12.8, 23:59 (3 weeks)
- Late Submission: 2024.12.9 ~ 2024.12.15, 23:59
 - Score * 0.8
 - No more submissions after 2024.12.16

Problems Emerged

- Improper TTL handling in traceroute / ping
 - when to do TTL--?
 - when to send ICMP time exceeded?
- Fail to maintain ARP cache entries / send ARP requests
- Imperfect longest prefix match
- Compilation & project structure problem
 - check before submission

Get to work
as soon as possible!

Q&A

Good luck

Acknowledgement

- This project is based on the CS118 class project by Alexander Afanasyev, UCLA.